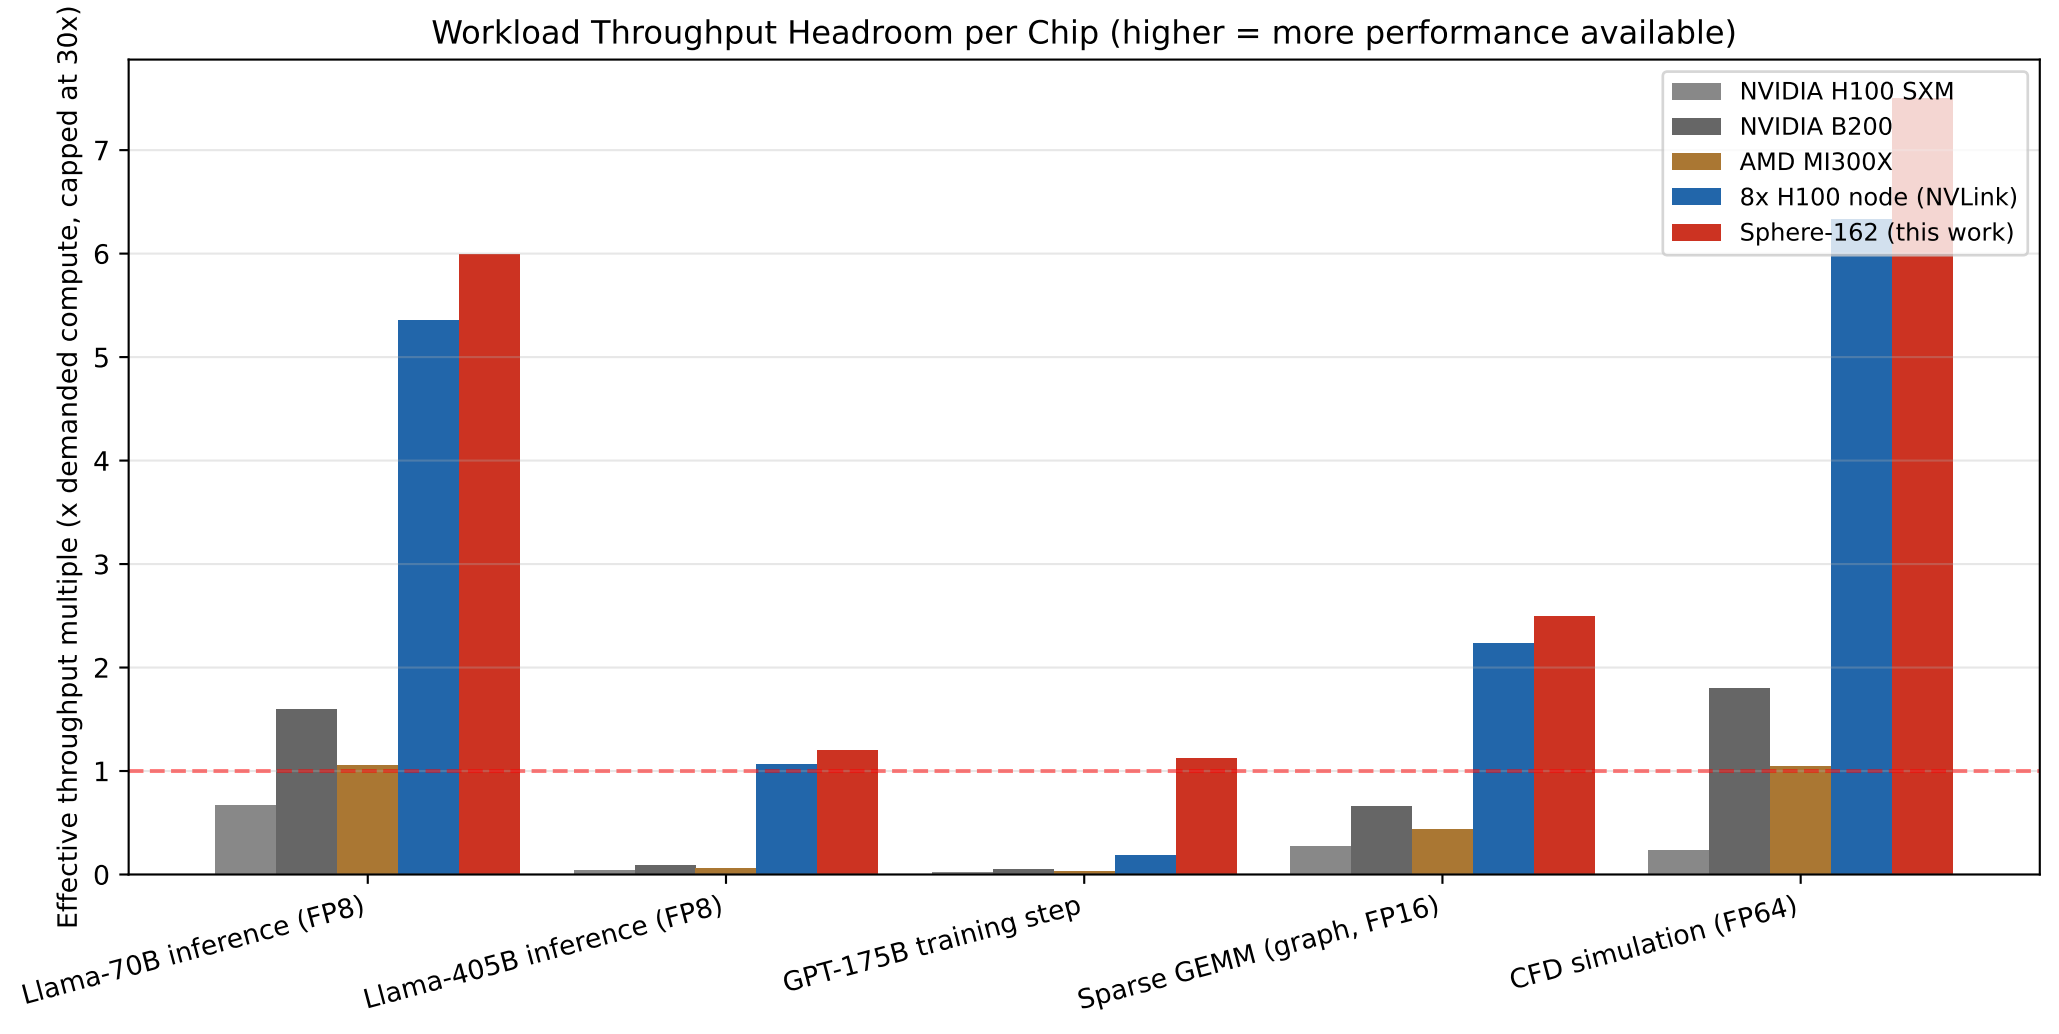
Workload Throughput Headroom per Chip (higher = more performance available)
Effective throughput multiple (x demanded compute, capped at 30x)
0
1
2
3
4
5
6
7
NVIDIA H100 SXM
NVIDIA B200
AMD MI300X
8x H100 node (NVLink)
Sphere-162 (this work)
Llama-70B inference (FP8)
Llama-405B inference (FP8)
GPT-175B training step
Sparse GEMM (graph, FP16)
CFD simulation (FP64)